

Event-driven architecture

Unpacking Knative



Jeff Barnes
GC Micro-mission

Event-driven (Serverless) Architecture

Event-Driven, sometimes also called serverless or functions as a service, is a computing execution model in which the infrastructure/provider dynamically manages the allocation of machine resources

- Code is deployed to a CSP or elsewhere
- Instance does not exist until invoked, spins up and scales as required, scales to zero when complete
- Changes IT requirements;
 - Security
 - Administration (patching, deployment)
 - Architecture
 - Cold start
- Currently various standards and little interoperability between CSP's
- Specific use cases include 'intents' for voice activated assistants (AWS Echo skills deployed in AWS Lambda)



IBM
OpenWhisk



AWS
Lambda



Google
Cloud Functions



Azure
Functions

.....

Event-driven (Serverless) Architecture

Originally found in public clouds, CSP instances are heading on-prem via appliance (Google Cloud Services Platform, Azure Stack, AWS Greengrass). A number of open-source versions are now available which can be deployed on-prem or in public cloud.

Examples include knative, OpenFaaS and Kubeless. This deck unpacks Knative, which can be used to auto-scale serverless-style functions, applications, and containers on Kubernetes



Knative



OpenFaaS



Kubeless

Initially serverless, Functions as a service (FaaS) and event-driven architecture referred to a microservice that is run 'on public cloud' compute engine only when invoked. Recently, serverless is also being used to describe managed container service (CaaS) and application deployments where the user is not responsible for the compute where/when the container runs (Azure AKS and ACI, AWS EKS and Fargate, Google GKE, Google App Engine, Ibm Kubernetes Service....)

Knative

Knative can be used to deploy serverless-style functions, applications, and containers to an auto-scaling runtime on Kubernetes. Also has many other features: deploy multiple versions, perform custom tasks on your application's source code, build reusable templates...

- [Install Knative on an Istio Cluster](#) (Istio injection enabled)
- [Select and install demo app](#) ([helloworld-go](#))

Invoke the app

- Determine if running 'LoadBalancer' or 'NodePort'

```
export IP_ADDRESS=$(kubectl get svc istio-ingressgateway --namespace istio-system --output 'jsonpath={.status.loadBalancer.ingress[0].ip}')
```

OR

```
export IP_ADDRESS=$(kubectl get node --output 'jsonpath={.items[0].status.addresses[0].address}') : $(kubectl get svc istio-ingressgateway --namespace istio-system --output 'jsonpath={.spec.ports[?(@.port==80)].nodePort}')
```


```
export HOST_URL=$(kubectl get ksvc helloworld-go --output jsonpath='{.status.domain}')
```

```
curl -H "Host: ${HOST_URL}" http://${IP_ADDRESS}
Hello World: Go Sample v1!
```

Knative – Cold Start

```
~/knative$ curl --header "Host: ${HOST_URL}" --write-out "  
> lookup           %{time_namelookup}  
> connect           %{time_connect}  
> appconnect        %{time_appconnect}  
> pretransfer       %{time_pretransfer}  
> redirect          %{time_redirect}  
> starttransfer     %{time_starttransfer}  
> total             %{time_total}\n" \  
> http://${IP_ADDRESS}  
Hello Go Sample v1!
```

lookup	0.000067
connect	0.002420
appconnect	0.000000
pretransfer	0.002490
redirect	0.000000
starttransfer	15.767278
total	15.767717



Cold Start time 15 sec (need to fix it in my demo)

← → ↻ ⓘ Not secure | 10.102.7.108

Hello Go Sample v1!

```
~/knative$ curl --header "Host: ${HOST_URL}" --write-out "  
lookup           %{time_namelookup}  
connect           %{time_connect}  
appconnect        %{time_appconnect}  
pretransfer       %{time_pretransfer}  
redirect          %{time_redirect}  
starttransfer     %{time_starttransfer}  
total             %{time_total}\n" http://${IP_ADDRESS}  
Hello Go Sample v1!  
  
lookup           0.000067  
connect           0.002420  
appconnect        0.000000  
pretransfer       0.002490  
redirect          0.000000  
starttransfer     15.767248  
total             15.767717  
~/knative$
```

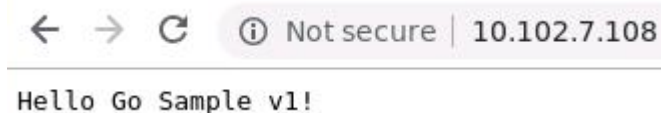
Knative – Subsequent Start

```
~/knative$ curl --header "Host: ${HOST_URL}" --write-out "
```

```
lookup           %{time_namelookup}  
connect          %{time_connect}  
appconnect       %{time_appconnect}  
pretransfer      %{time_pretransfer}  
redirect         %{time_redirect}  
starttransfer    %{time_starttransfer}  
total            %{time_total}\n" http://${IP_ADDRESS}  
Hello Go Sample v1!
```

```
lookup           0.000073  
connect          0.003975  
appconnect       0.000000  
pretransfer      0.004061  
redirect         0.000000  
starttransfer    0.062721  
total            0.062793
```

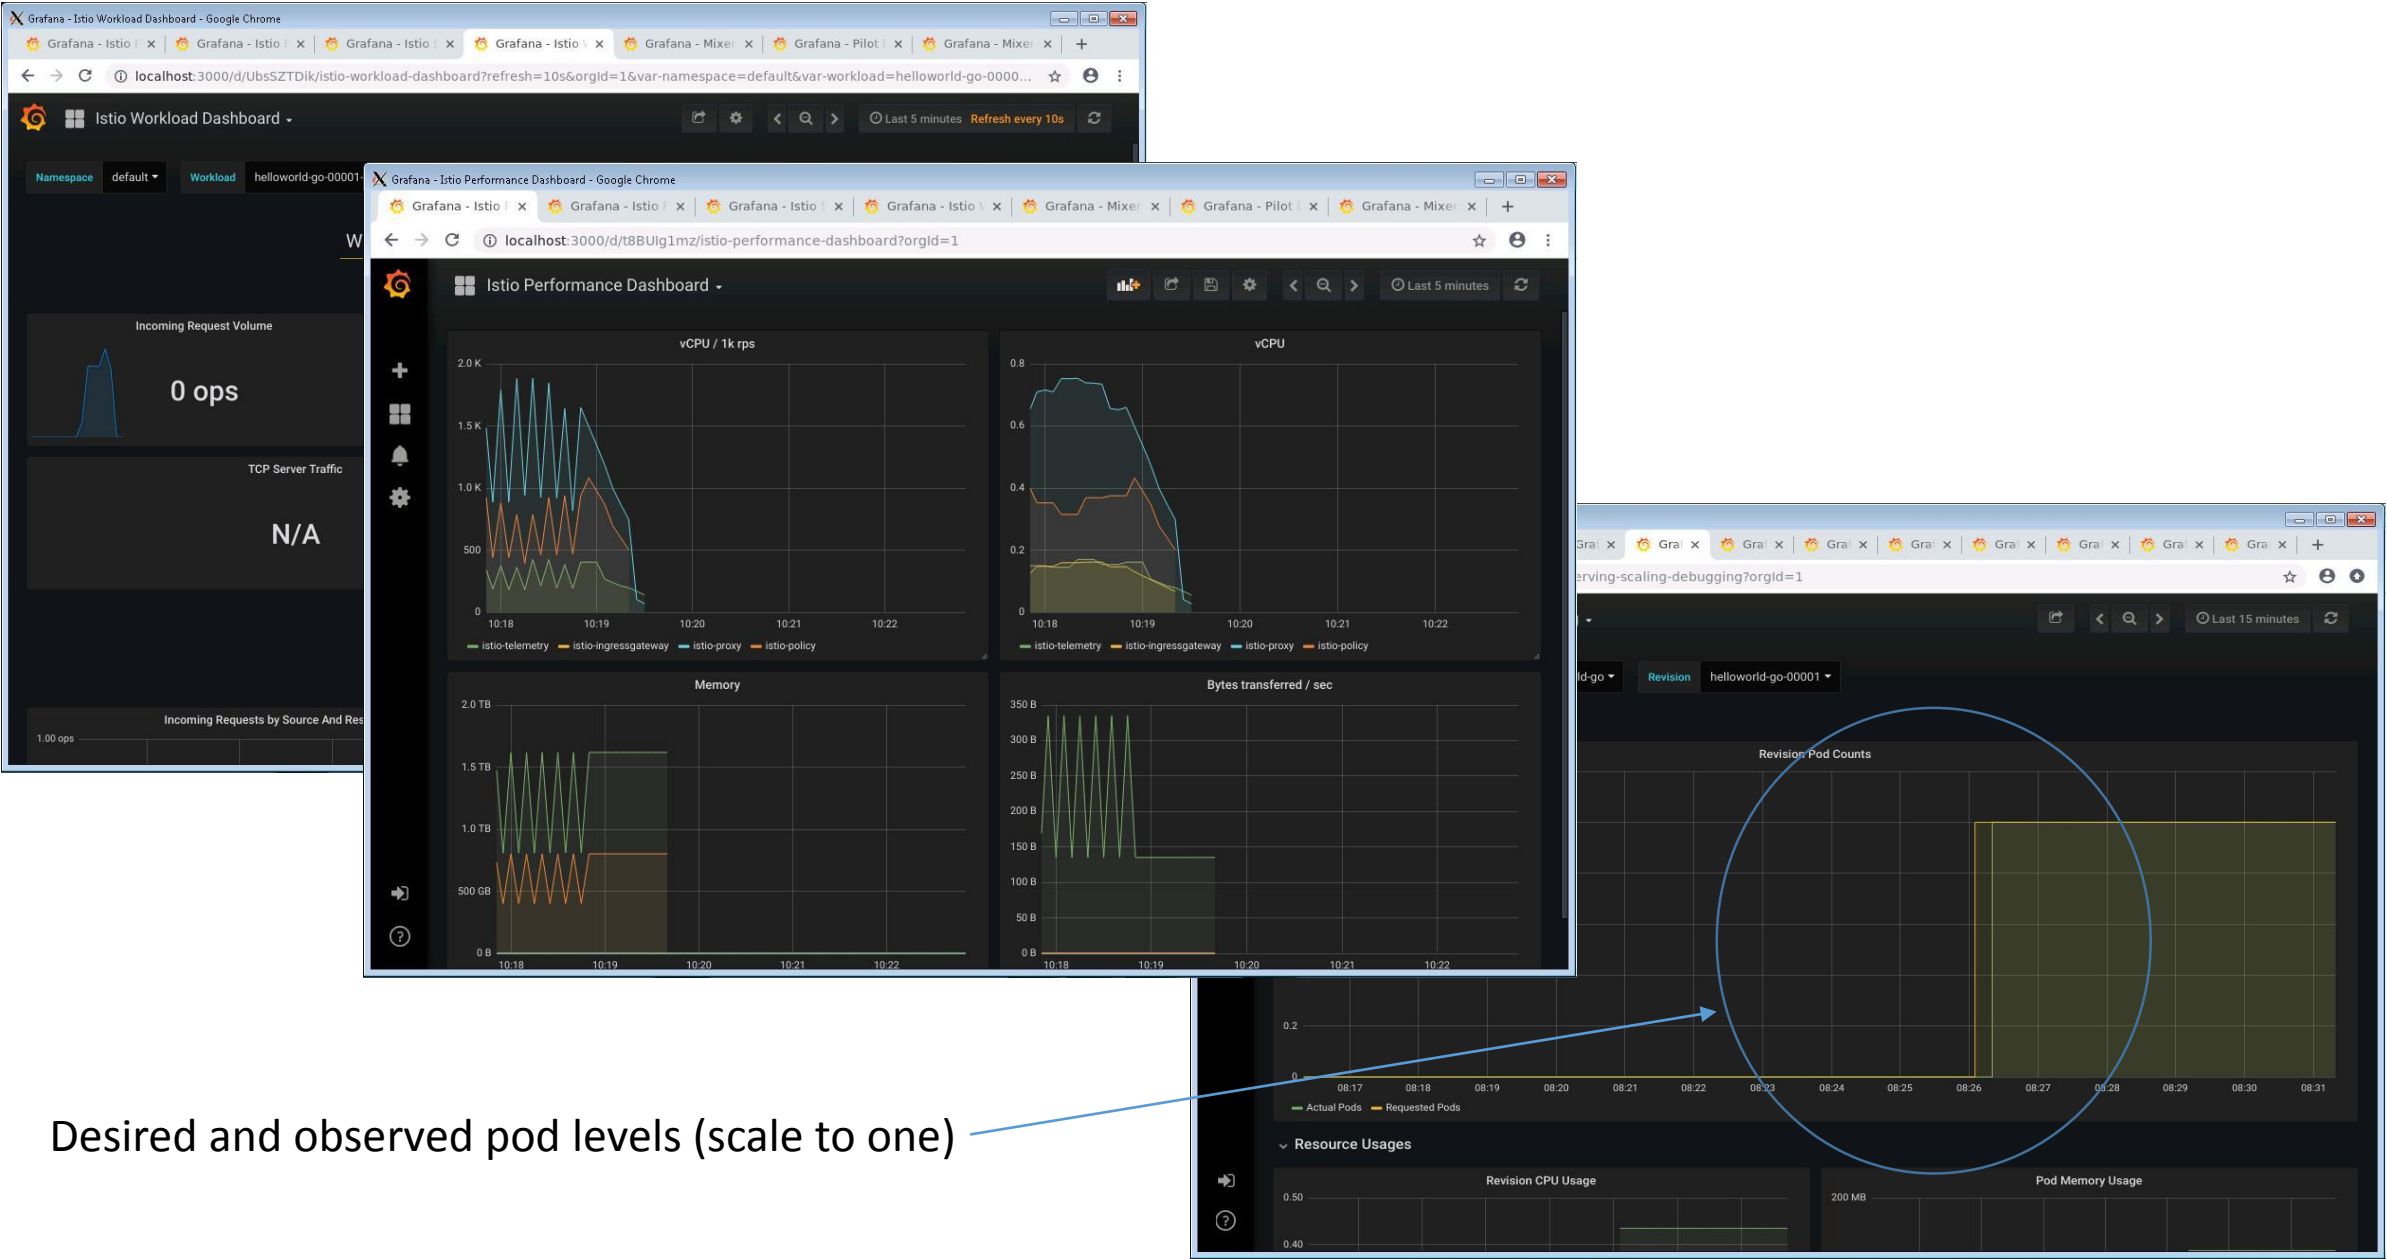
Subsequent response time < 1s



← → ↻ ⓘ Not secure | 10.102.7.108
Hello Go Sample v1!

```
~/knative$ curl --header "Host: ${HOST_URL}" --write-out "  
lookup           %{time_namelookup}  
connect          %{time_connect}  
appconnect       %{time_appconnect}  
pretransfer      %{time_pretransfer}  
redirect         %{time_redirect}  
starttransfer    %{time_starttransfer}  
total            %{time_total}\n" http://${IP_ADDRESS}  
Hello Go Sample v1!  
  
lookup           0.000073  
connect          0.003975  
appconnect       0.000000  
pretransfer      0.004061  
redirect         0.000000  
starttransfer    0.062721  
total            0.062793  
~/knative$
```

Knative – Performance Monitoring Grafana



Knative – Scale based on workload


[hey](#) runs number of requests in the provided concurrency level and prints stats.

- n Number of requests to run. Default is 200.
- c Number of requests to run concurrently. Total number of requests cannot be smaller than the concurrency level. Default is 50.

```
~/go/bin$ ./hey -host helloworld-go.default.example.com -c 500 -n 150000 "http://${IP_ADDRESS}?sleep=1000"
```

Summary:

Total:	291.8676 secs
Slowest:	17.4429 secs
Fastest:	0.0025 secs
Average:	0.9101 secs
Requests/sec:	513.9316
Total data: 3002220 bytes	
Size/request: 20 bytes	



Send 150,000 requests (with 500 requests in parallel), each taking 1 second

Latency distribution:

10%	in 0.0041 secs
25%	in 0.0050 secs
50%	in 0.0083 secs
75%	in 1.3462 secs
90%	in 3.4178 secs
95%	in 4.5164 secs
99%	in 6.8141 secs

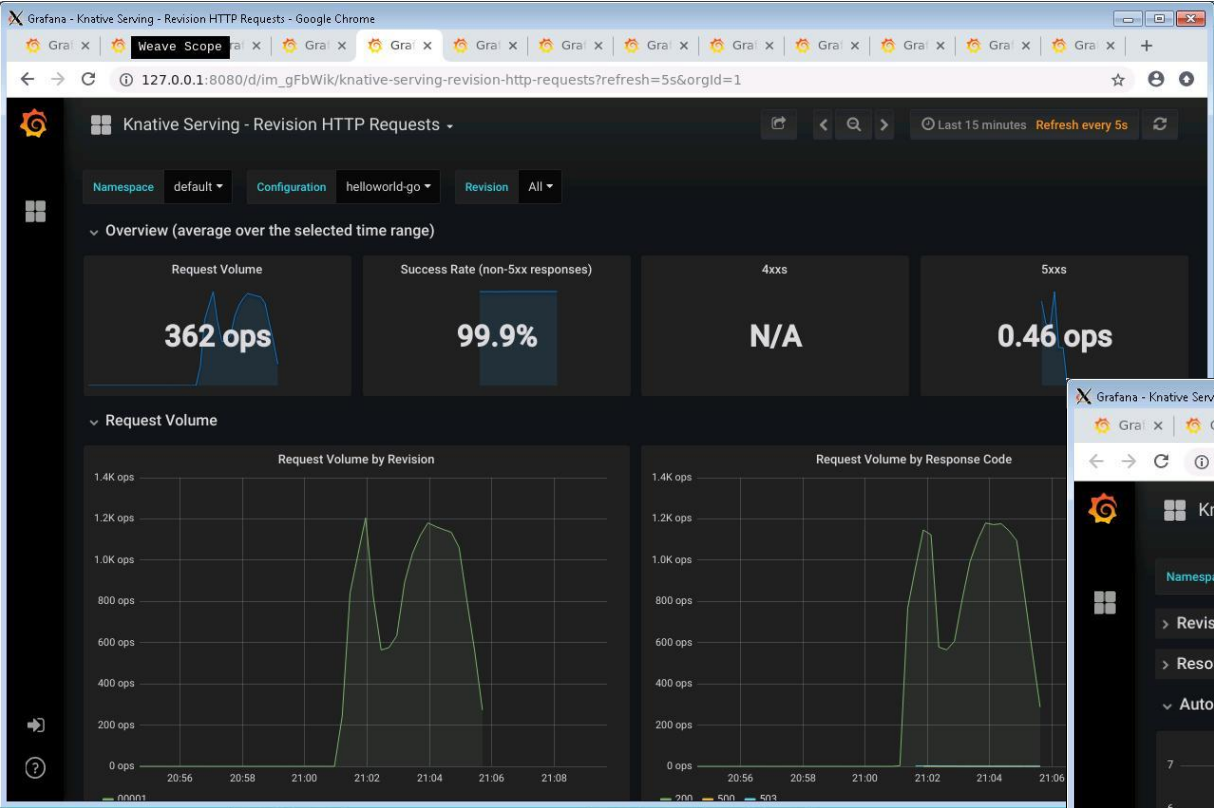
Details (average, fastest, slowest):

DNS+dialup:	0.0002 secs,	0.0025 secs,	17.4429 secs
DNS-lookup:	0.0000 secs,	0.0000 secs,	0.0000 secs
req write:	0.0002 secs,	0.0000 secs,	0.1051 secs
resp wait:	0.9091 secs,	0.0024 secs,	17.4428 secs
resp read:	0.0005 secs,	0.0000 secs,	0.1377 secs

Status code distribution:

[200]	149940 responses
[503]	60 responses

Knative – Scale based on workload w/Performance Monitoring Grafana



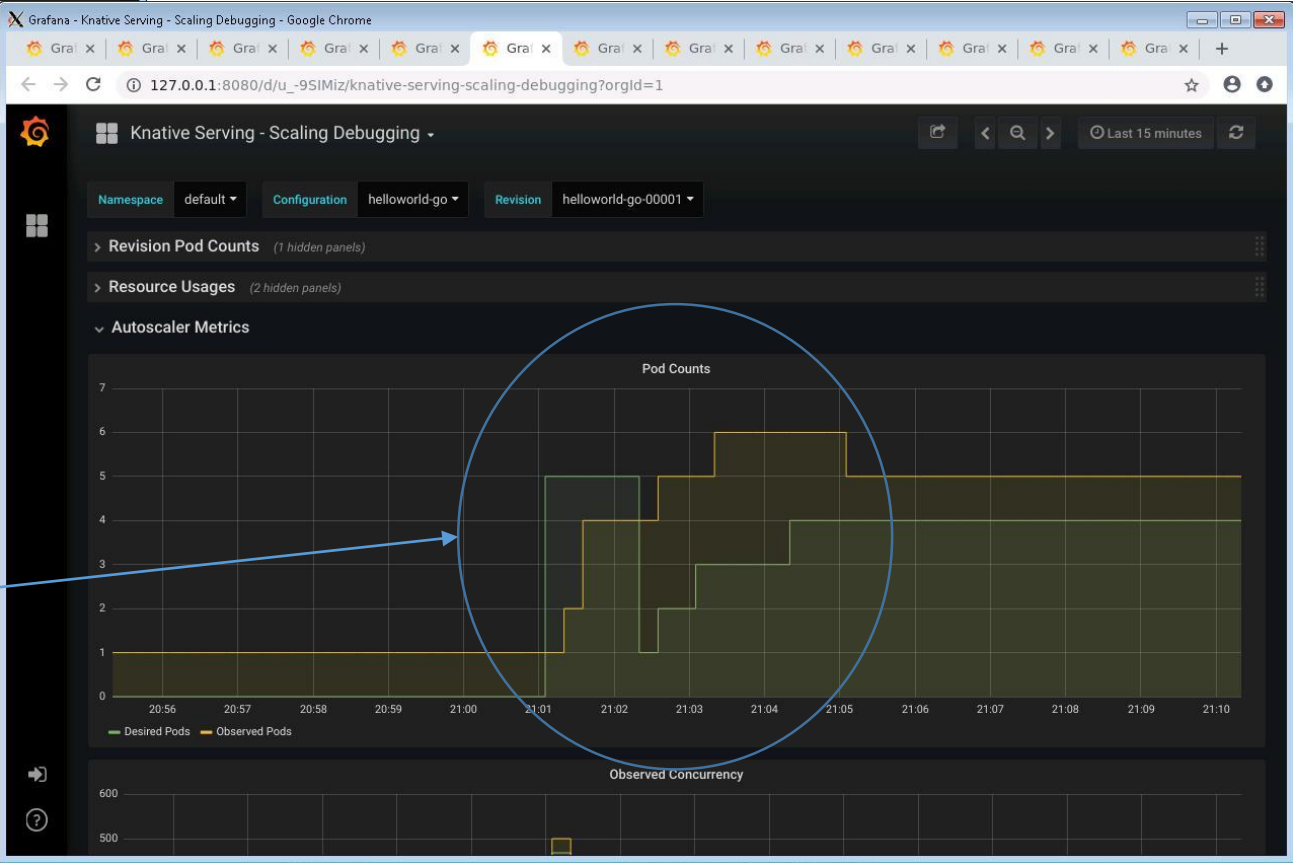
← → ↻ ⓘ Not secure | 10.102.7.108

Hello Go : ← → ↻ ⓘ Not secure | 10.102.7.108

Hello Go Samp ← → ↻ ⓘ Not secure | 10.102.7.108

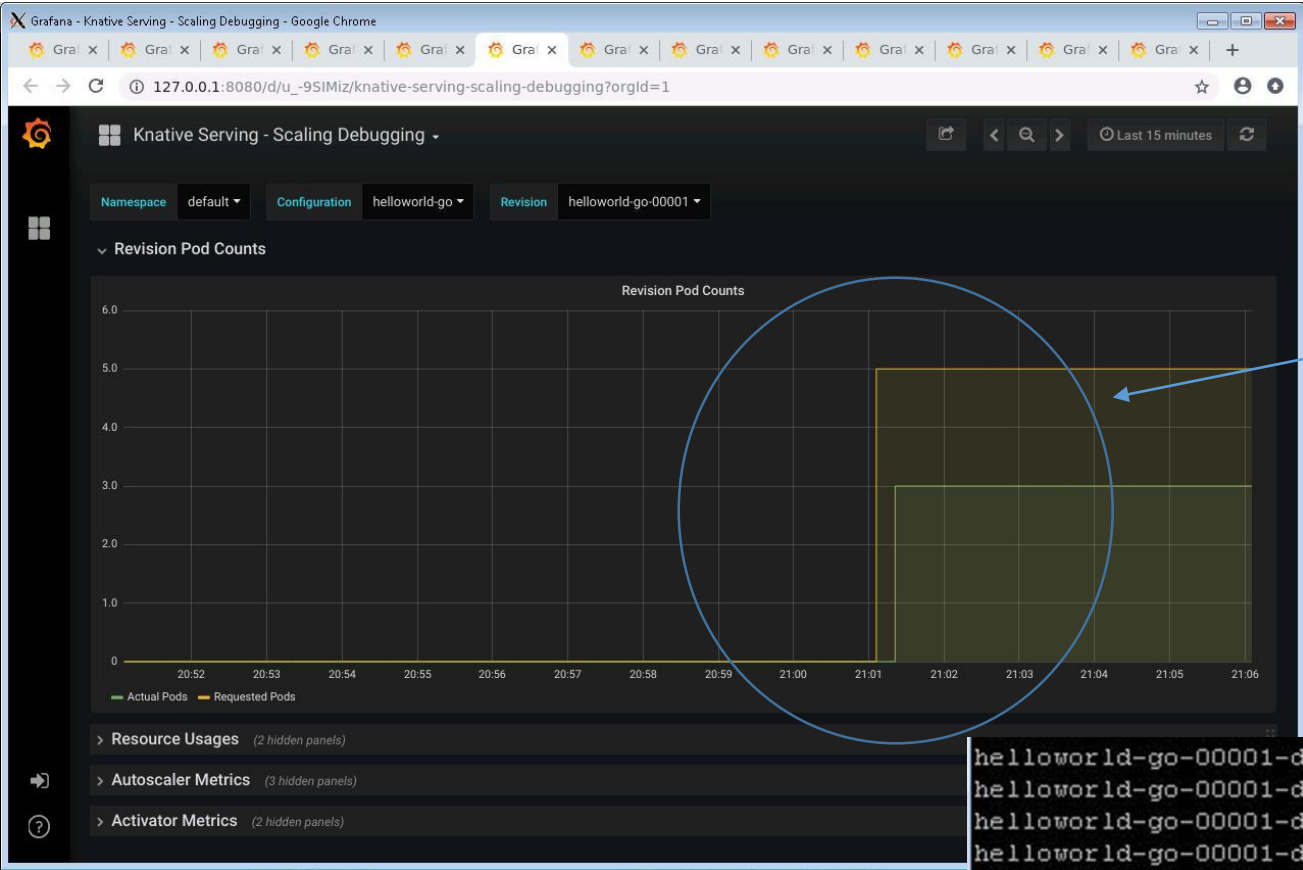
Hello Go Sample v1!

Desired and observed pod levels



Knative – Scale based on workload

Knative Serving, by default, has a concurrent requests target of 100. Sending 500 concurrent requests causes autoscaling to note that it needs to run 5 Pods to satisfy this level



Desired and observed pod levels

Running and terminating Pods
(as required)

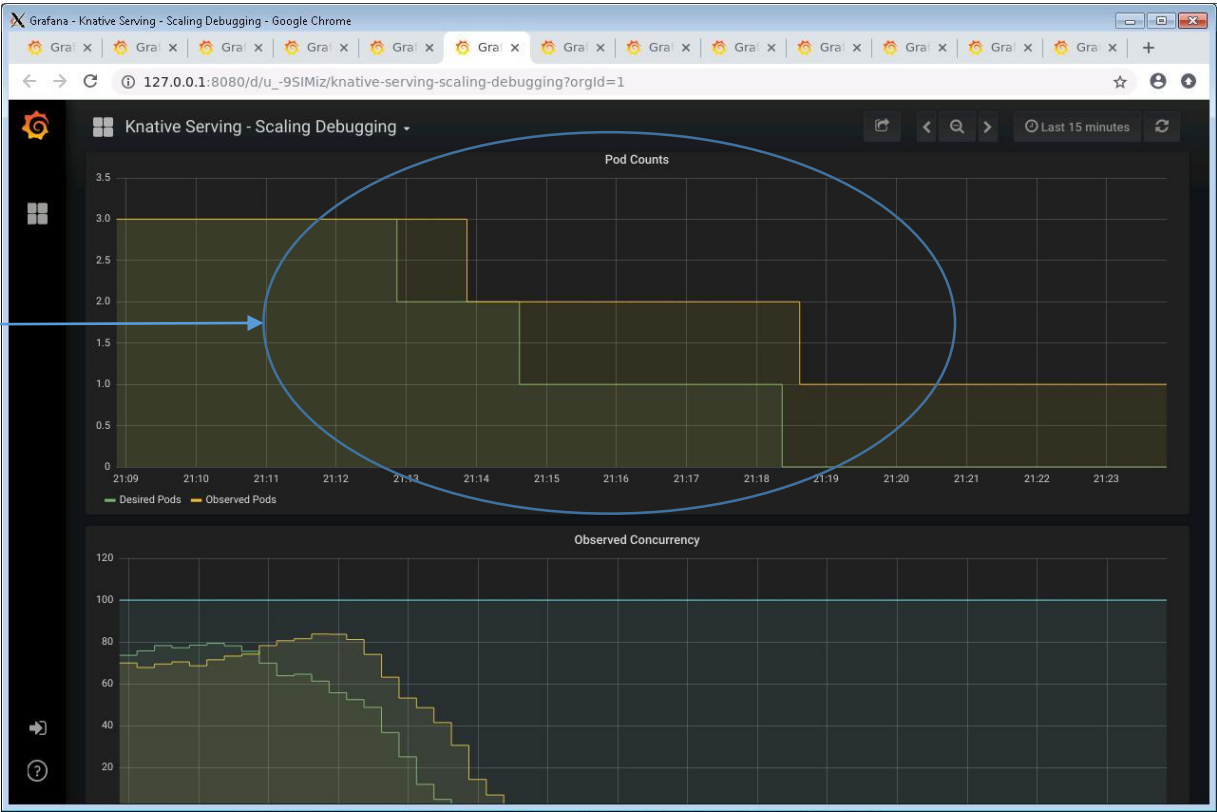
helloworld-go-00001-deployment-94667496b-6tccn	3/3	Running	0	52s
helloworld-go-00001-deployment-94667496b-9lqpj	3/3	Running	0	2m48s
helloworld-go-00001-deployment-94667496b-h8s8c	0/3	Terminating	0	2m48s
helloworld-go-00001-deployment-94667496b-l2xgv	2/3	Terminating	0	2m48s
helloworld-go-00001-deployment-94667496b-m14cw	2/3	Terminating	0	2m46s
helloworld-go-00001-deployment-94667496b-r15tr	2/3	Terminating	0	2m48s
helloworld-go-00001-deployment-94667496b-xrzjm	3/3	Running	0	84s

Knative – Scale based on workload w/Performance Monitoring Grafana

helloworld-go-00001-deployment-94667496b-2npf2	0/3	Terminating	0	7m20s
helloworld-go-00001-deployment-94667496b-6tccn	3/3	Running	0	8m46s
helloworld-go-00001-deployment-94667496b-xrzjm	3/3	Running	0	9m6s

Running and terminating pods as traffic decreases

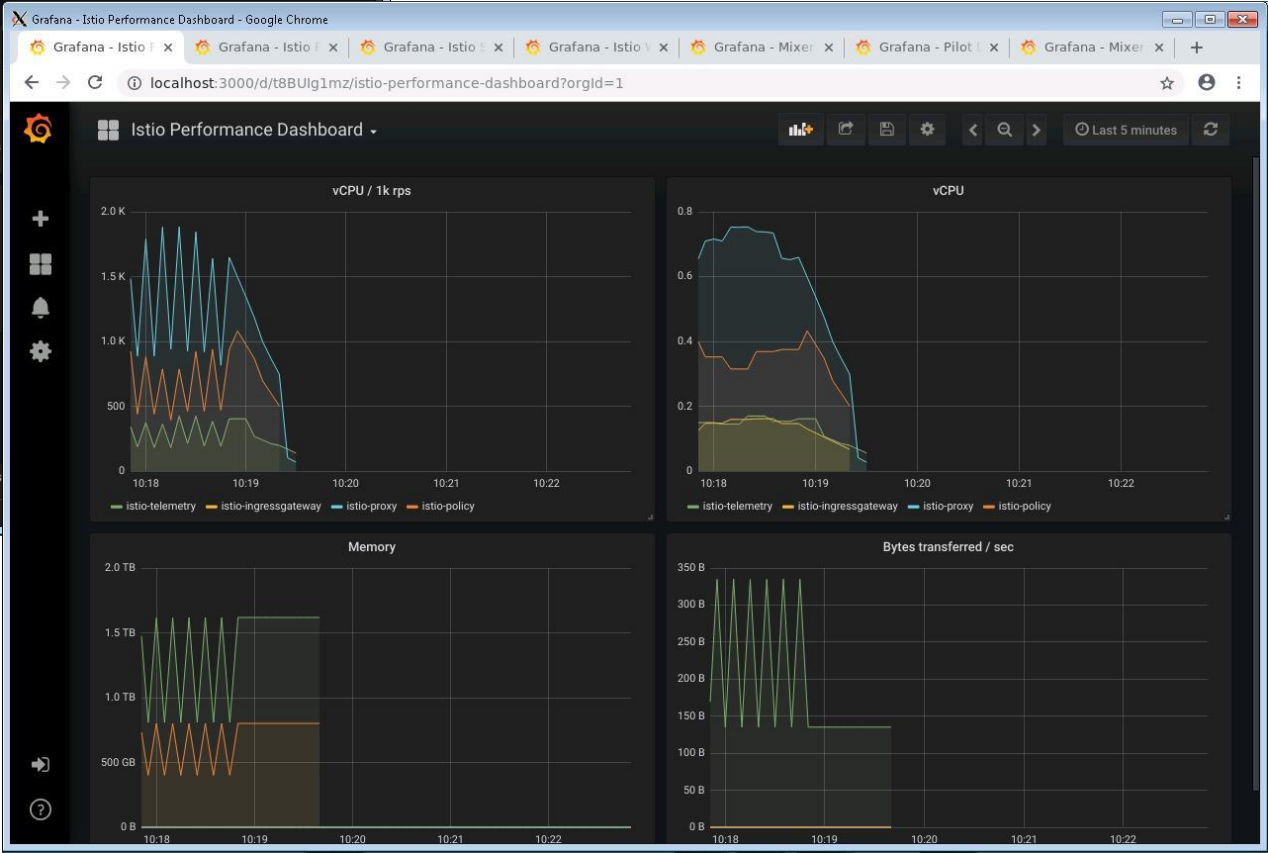
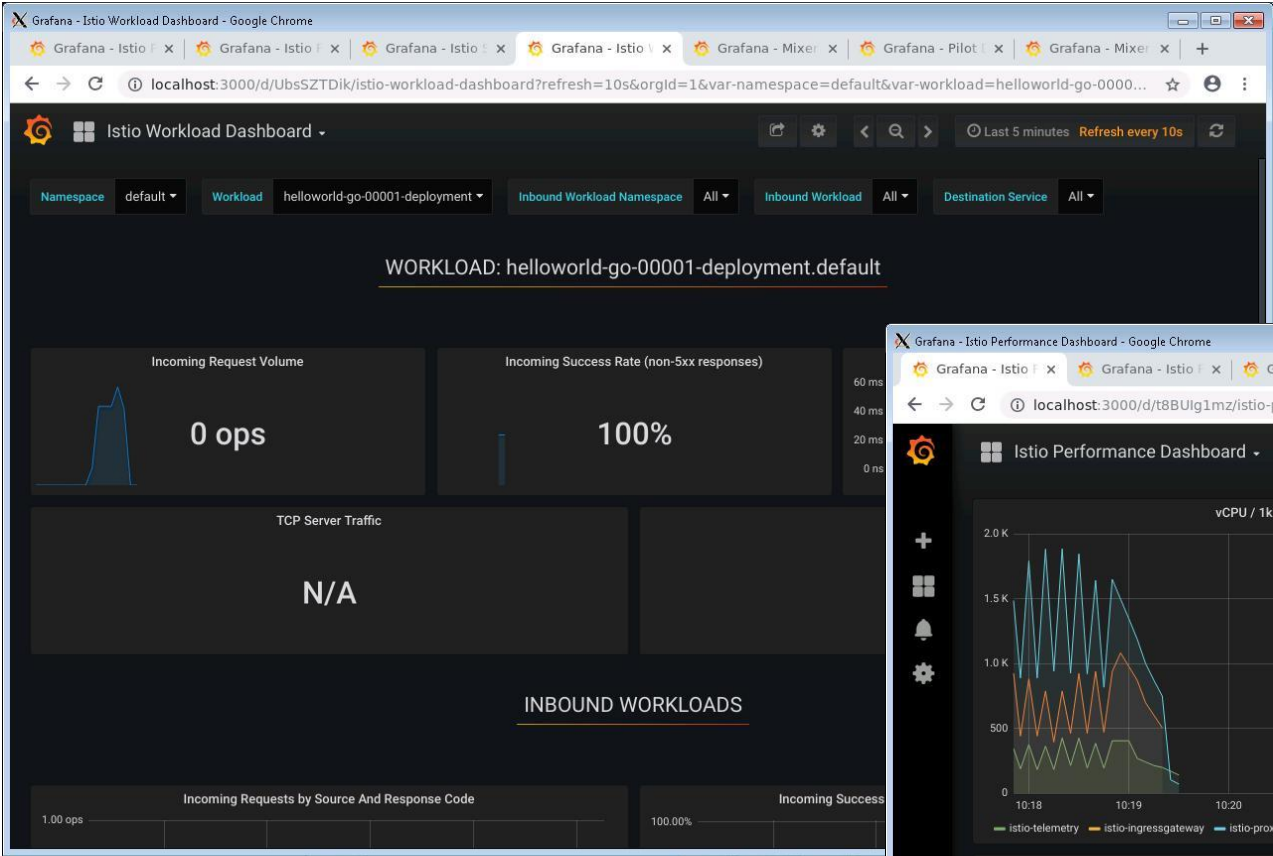
Desired and observed pod levels as traffic decreases



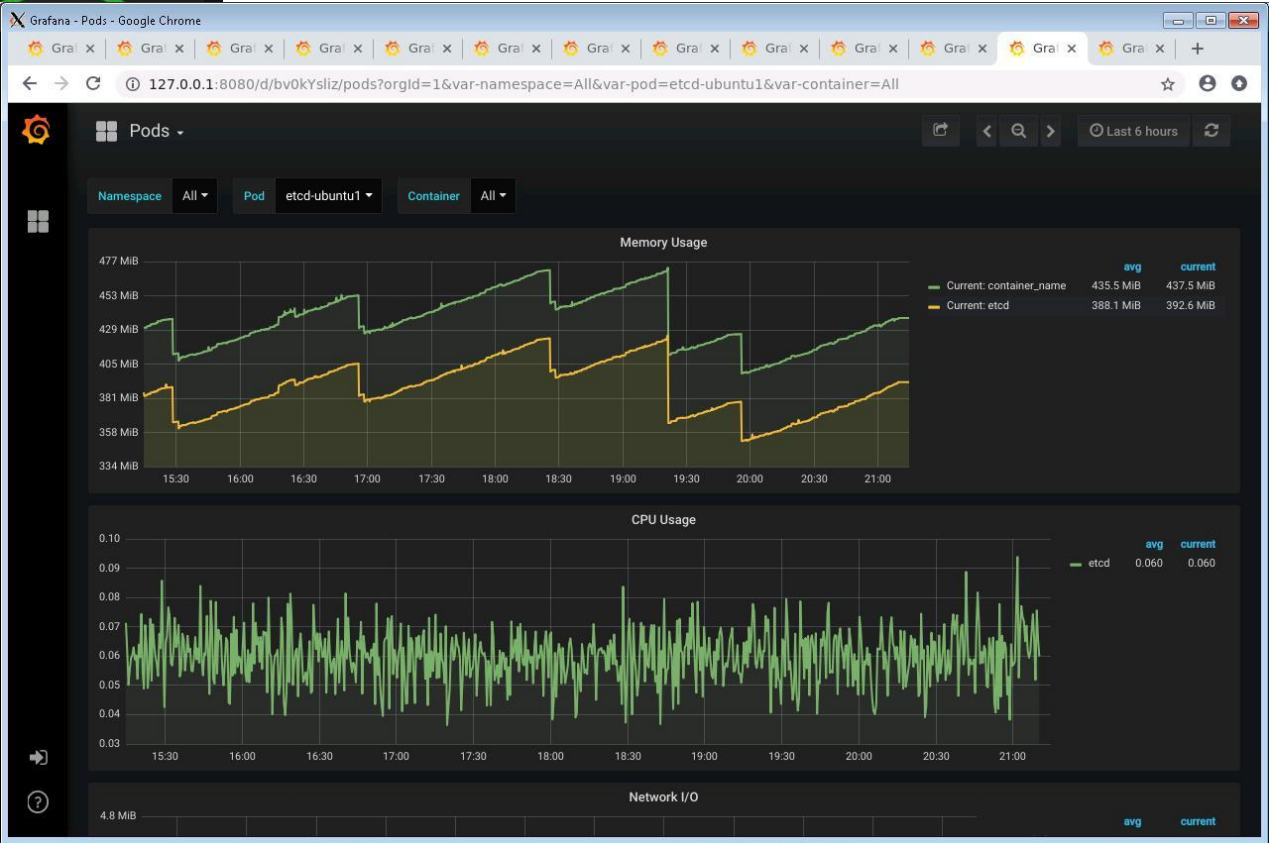
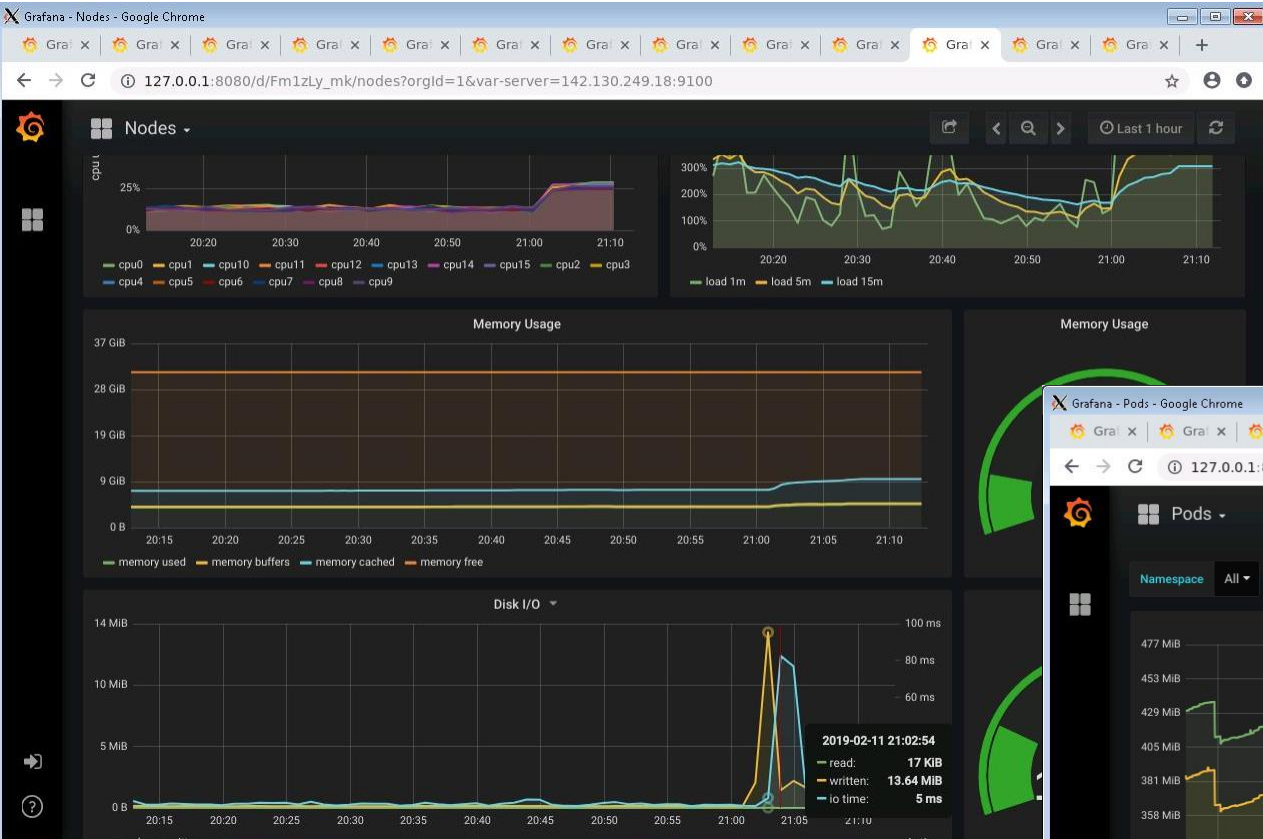
Desired and observed concurrency

helloworld-go-00001-deployment-94667496b-xrzjm	2/3	Terminating	0	18m
--	-----	-------------	---	-----

Knative – Performance Monitoring Grafana



Knative – Performance Monitoring Grafana



Knative – Demo

\$ kubectl get pods

Runs as required

helloworld-go-00001-deployment-94667496b-68x2q	0/3	Init:0/1	0	7s
helloworld-go-00001-deployment-94667496b-9p2cq	0/3	Init:0/1	0	5s
helloworld-go-00001-deployment-94667496b-9sff9	0/3	PodInitializing	0	7s
helloworld-go-00001-deployment-94667496b-jqbjp	0/3	PodInitializing	0	7s
helloworld-go-00001-deployment-94667496b-pb2zq	0/3	Init:0/1	0	7s
helloworld-go-00001-deployment-94667496b-68x2q	0/3	Init:0/1	0	11s
helloworld-go-00001-deployment-94667496b-9p2cq	0/3	PodInitializing	0	9s
helloworld-go-00001-deployment-94667496b-9sff9	2/3	Running	0	11s
helloworld-go-00001-deployment-94667496b-jqbjp	0/3	PodInitializing	0	11s
helloworld-go-00001-deployment-94667496b-pb2zq	0/3	Init:0/1	0	11s

Then self-terminates when no longer required

helloworld-go-00001-deployment-94667496b-68x2q	2/3	Terminating	0	95s
helloworld-go-00001-deployment-94667496b-9p2cq	2/3	Terminating	0	93s
helloworld-go-00001-deployment-94667496b-9sff9	2/3	Terminating	0	95s
helloworld-go-00001-deployment-94667496b-jqbjp	3/3	Running	0	95s
helloworld-go-00001-deployment-94667496b-pb2zq	2/3	Terminating	0	95s
helloworld-go-00001-deployment-94667496b-68x2q	2/3	Terminating	0	2m47s
helloworld-go-00001-deployment-94667496b-9p2cq	2/3	Terminating	0	2m45s
helloworld-go-00001-deployment-94667496b-9sff9	2/3	Terminating	0	2m47s
helloworld-go-00001-deployment-94667496b-jqbjp	2/3	Terminating	0	2m47s
helloworld-go-00001-deployment-94667496b-pb2zq	2/3	Terminating	0	2m47s

Knative – Demo