

# Service Mesh

Unpacking Istio



Jeff Barnes  
GC Micro-mission

# Unscripted post-demo deck

Provide background for the unscripted Service Mesh walkthrough

- How we got here and why
- Where to next

Service Mesh aligns to a number of current deliverables and overall guidance/standards in the GC, SSC and EA:

- Provides a mechanism to implement GC guidance such as and Directive on Management of Information Technology, with an updated Appendix C and D <https://www.tbs-sct.gc.ca/pol/doc-eng.aspx?id=15249> and the Digital Operations Strategic Plan: 2018-2022 <https://www.canada.ca/en/government/system/digital-government/digital-operations-strategic-plan-2018-2022.html>
- Use case for software defined overlays, micro segmentation and even zero-trust networking (related to the E2E and CBSA zoning, SCED and the CSE guidance updates)
- Supports partner application migration, especially monitoring
- Components support the Government of Canada Standards on APIs <https://www.canada.ca/en/government/publicservice/modernizing/government-canada-standards-apis.html>, API gateway and framework (including distributed tracing/open tracing)
- Open source, globally adopted
- Numerous other cloud-native benefits (container orchestration, reduces hardware footprint, declarative, application support including circuit breaking, fault tolerance,...)

# Origin Story

In an effort to better understand security for containers, orchestration, CI/CD; created a simple POC including Container Security Providers (Neuvector, AquaSec), CI/CD (Jenkins), Docker, minikube, Cloud Foundry, OpenShift; on bare metal and in various public clouds

Quickly discovered introductory tutorials often did not include security requirements and challenges concerning containers, especially in enterprise environments that must meet NIST SP800-53

Started with Go, Node and other runtimes; created and then secured Docker images, including the adoption of multi-stage and non-root builds, vulnerability scanning, code analysis, and other security recommendations\*

Subsequently led to a POC for securing platforms, especially Kubernetes and OpenShift OKD; which in turn included Helm for application deployment; and now where we are now – investigating Service Mesh using Istio

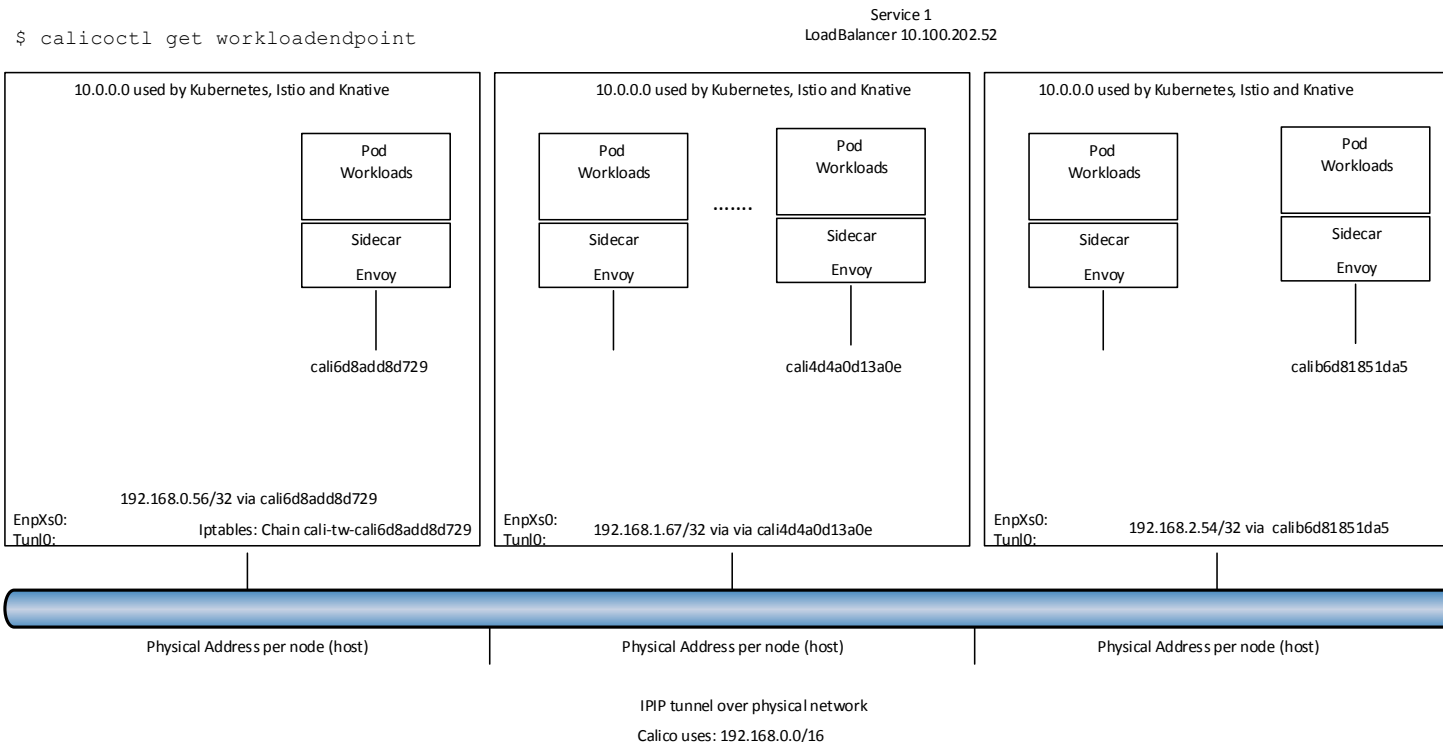


# Container Orchestrator Networking

Kubernetes 1.13 deployment, Calico IPIP + Kubeproxy (iptables) for CNI

High-level overview of Kubernetes networking.

- Physical Addresses assigned to hosts (nodes) that comprise the cluster
- Calico uses 192.168.0.0/16 for a pseudo-overlay on top of the nodes
- 10.0.0.0 used by Kubernetes (and Istio)

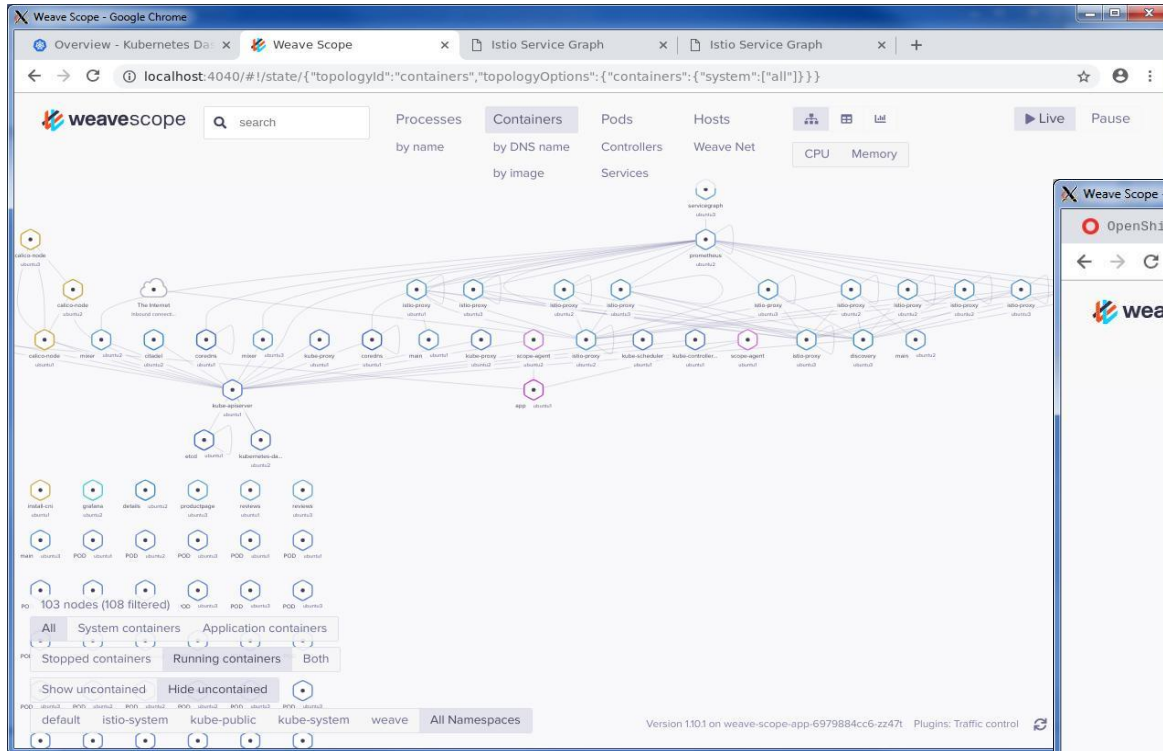


```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: my-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: db
  ingress:
    - from:
      - namespaceSelector:
          matchLabels:
            project: myproject
      - podSelector:
          matchLabels:
            role: frontend
    ports:
      - protocol: TCP
        port: 6379
```

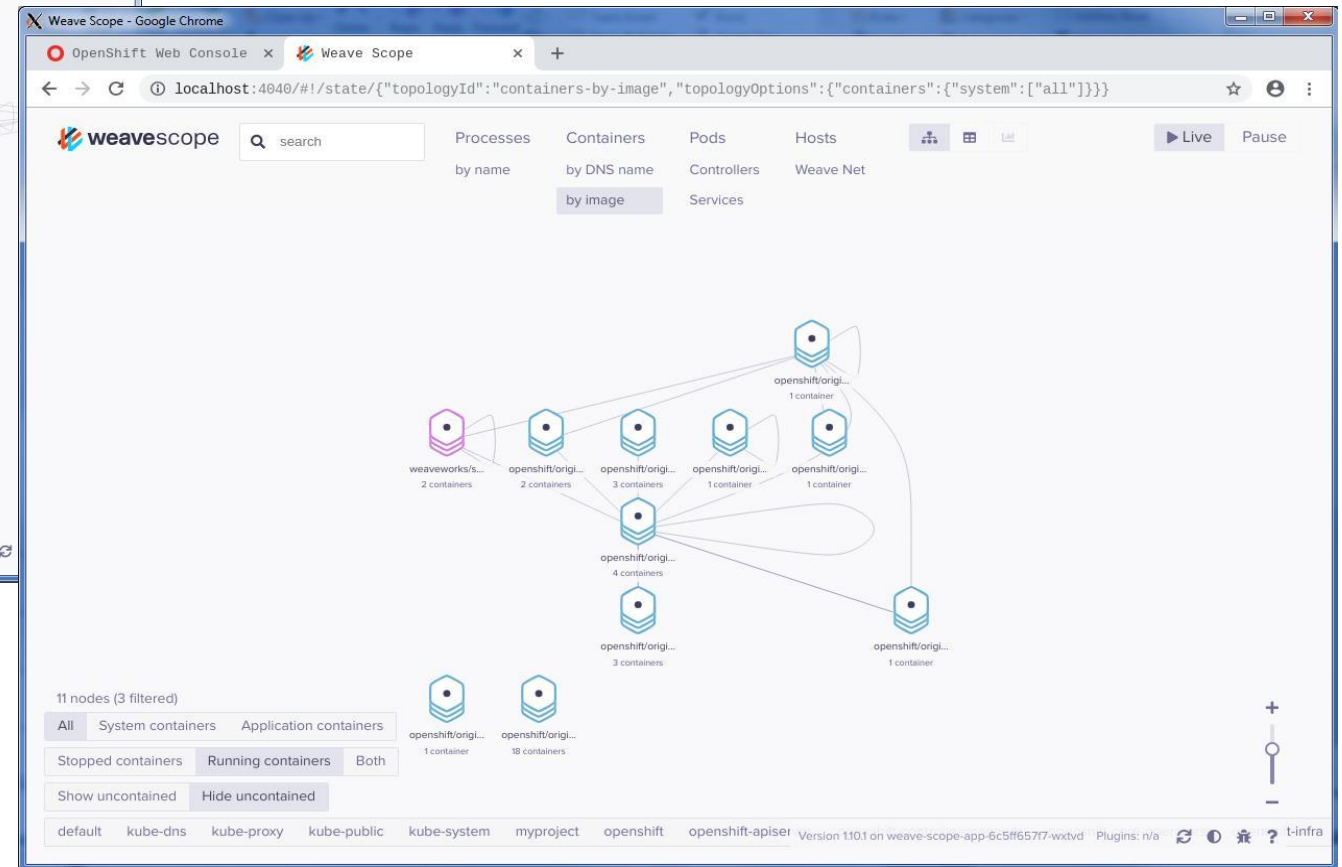
Example Network Policy

# Infrastructure Visualization

Using Weavescope to visual Kubernetes and OpenShift clusters



OpenShift cluster



# Why a Service Mesh

Decided to install a Service Mesh to better understand deployment, security and support of microservices and to examine use case for zoning, micro-segmentation, zero-trust networking, and cloud-native deployment

*“Much as Kubernetes abstracts the underlying compute capability from the hardware and operating system, so too do service meshes abstract routing and connectivity from the underlying networking. Once the routing and connectivity is separated from the infrastructure and can be treated as code, it becomes practical to have a truly application tailored, least privilege approach to microsegmentation.”*

Out of the box, Istio 1.0.5 comes with a number of features and components; including Prometheus and Grafana (logging/telemetry), Visualization, Zipkin and Jaeger for distributed tracing; plus all the benefits of sidecar proxies: Request Routing (for Canary or A/B deployments), Fault Injection, Traffic Shifting/control, Security/segmentation, mTLS, Health checking and on and on (all without changing any application code)

Envoy can stand alone and is included in many service meshes, such as Istio, AWS App Mesh....



Consul is a common proxy used for building service meshes in heterogeneous technical environments



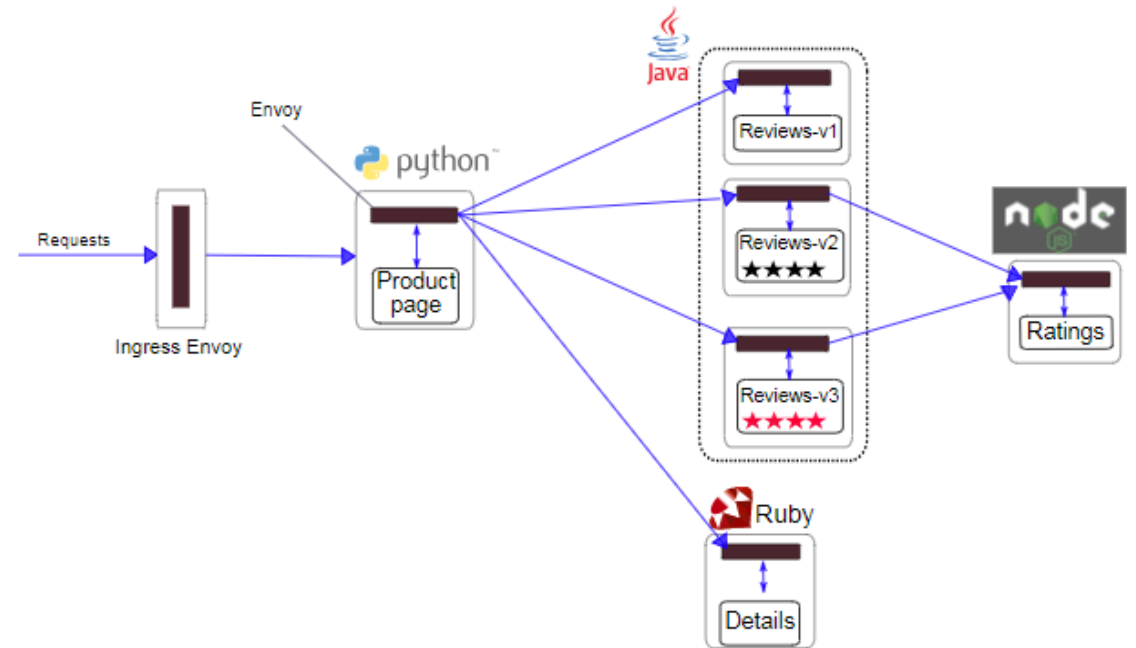
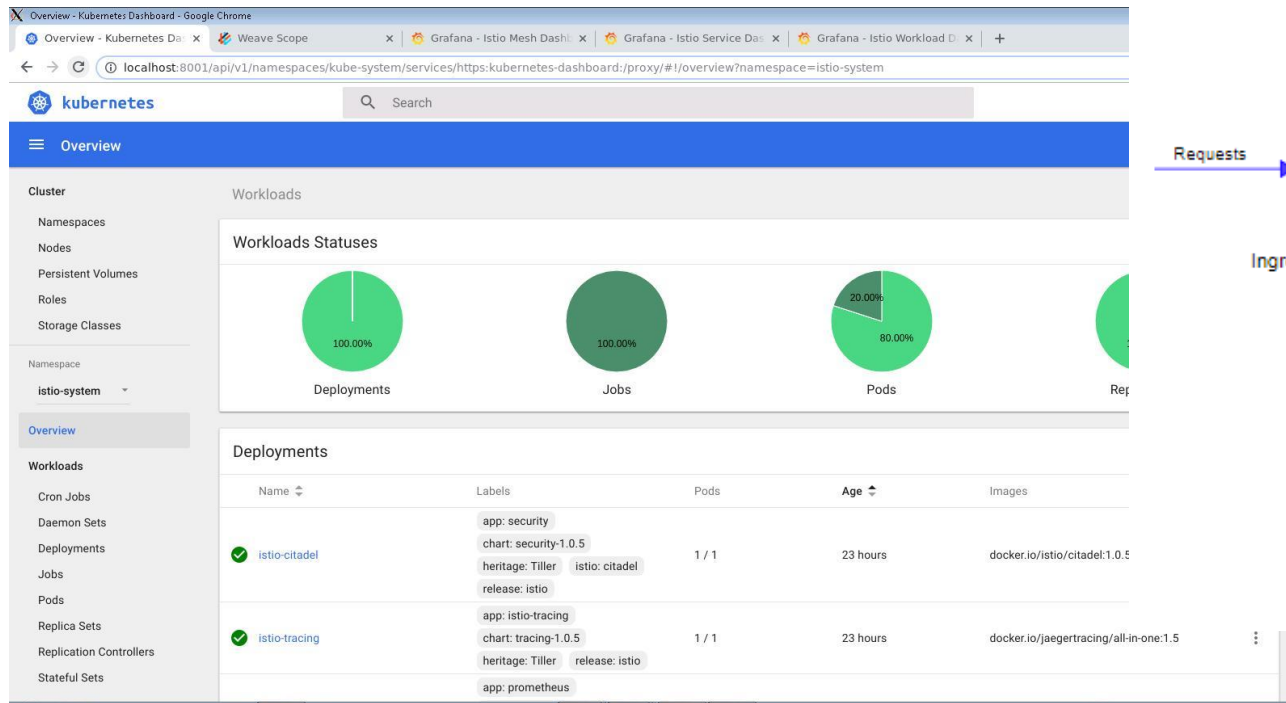
# Istio

Install in Istio-system namespace of a current kubernetes cluster. Main components include:

- Envoy - an open-source edge and service proxy designed for cloud-native applications
- Mixer - a platform-independent component responsible for enforcing access control and usage policies
- Pilot - provides service discovery for the Envoy sidecars and traffic management capabilities

We install two services in the default namespace

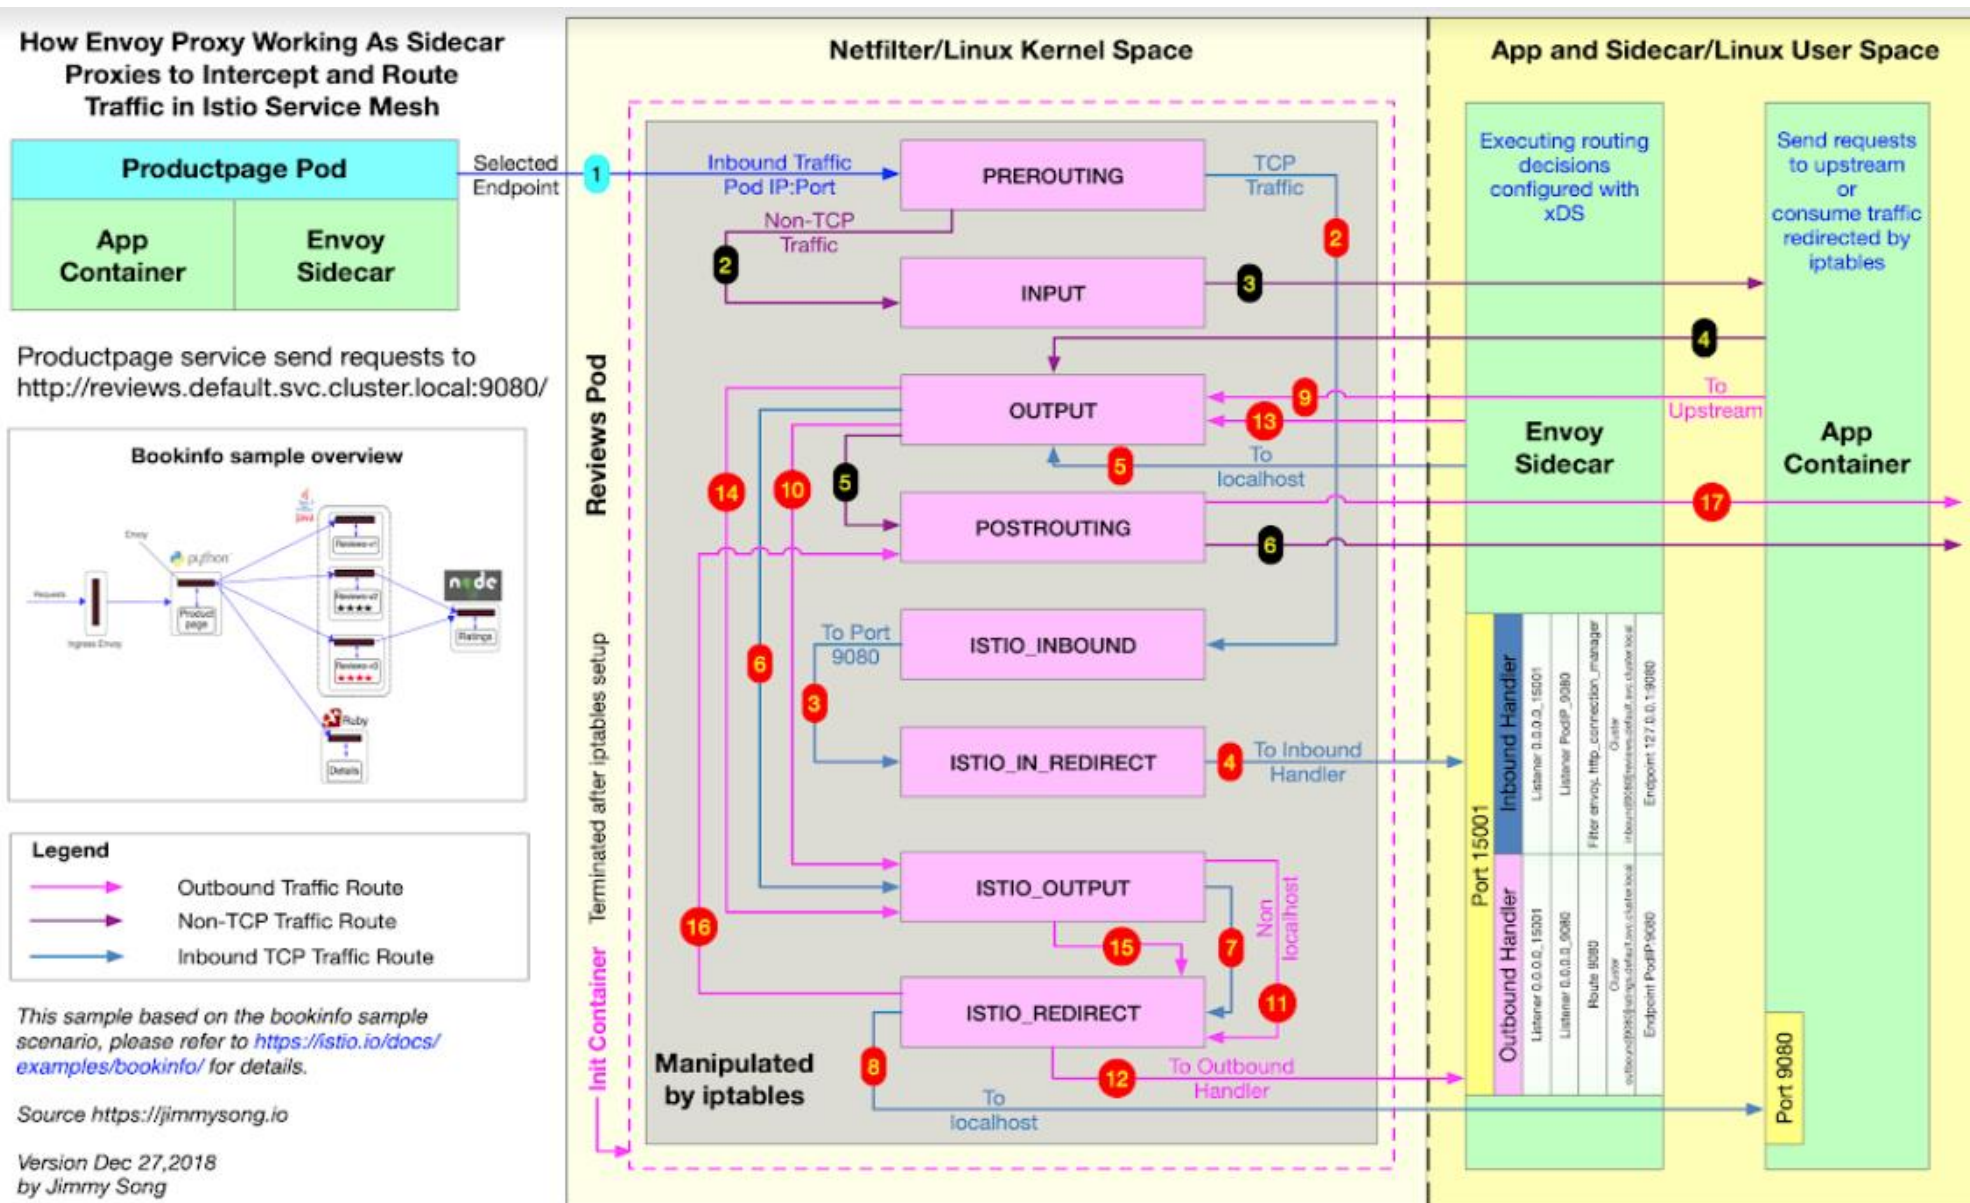
- **Main** – a 3 pod Simple web server
- **Bookinfo** – Example microservices for testing



Bookinfo Application



# Underneath the hood





# Traffic Shaping

Performed the following on **Bookinfo**, works as intended

[Request routing](#) This task will first direct all incoming traffic for the Bookinfo application to the v1 version of the reviews service. It will then send traffic only from a specific test user to version v2, leaving all other users unaffected.

[Fault injection](#) We will now use Istio to test the resiliency of the Bookinfo application by injecting an artificial delay in requests between the reviews:v2 and ratings services. Observing the resulting behavior as the test user, we will notice that the v2 version of the reviews service has a bug. Note that all other users are unaware of this testing against the live system.

[Traffic Shifting](#) Finally, we will use Istio to gradually migrate traffic for all users from to a v3 version of the reviews service, one which includes the fix for the bug discovered in v2.

[Collecting metrics](#) This task will configure Mixer to collect a uniform set of metrics across all services in the Bookinfo application.

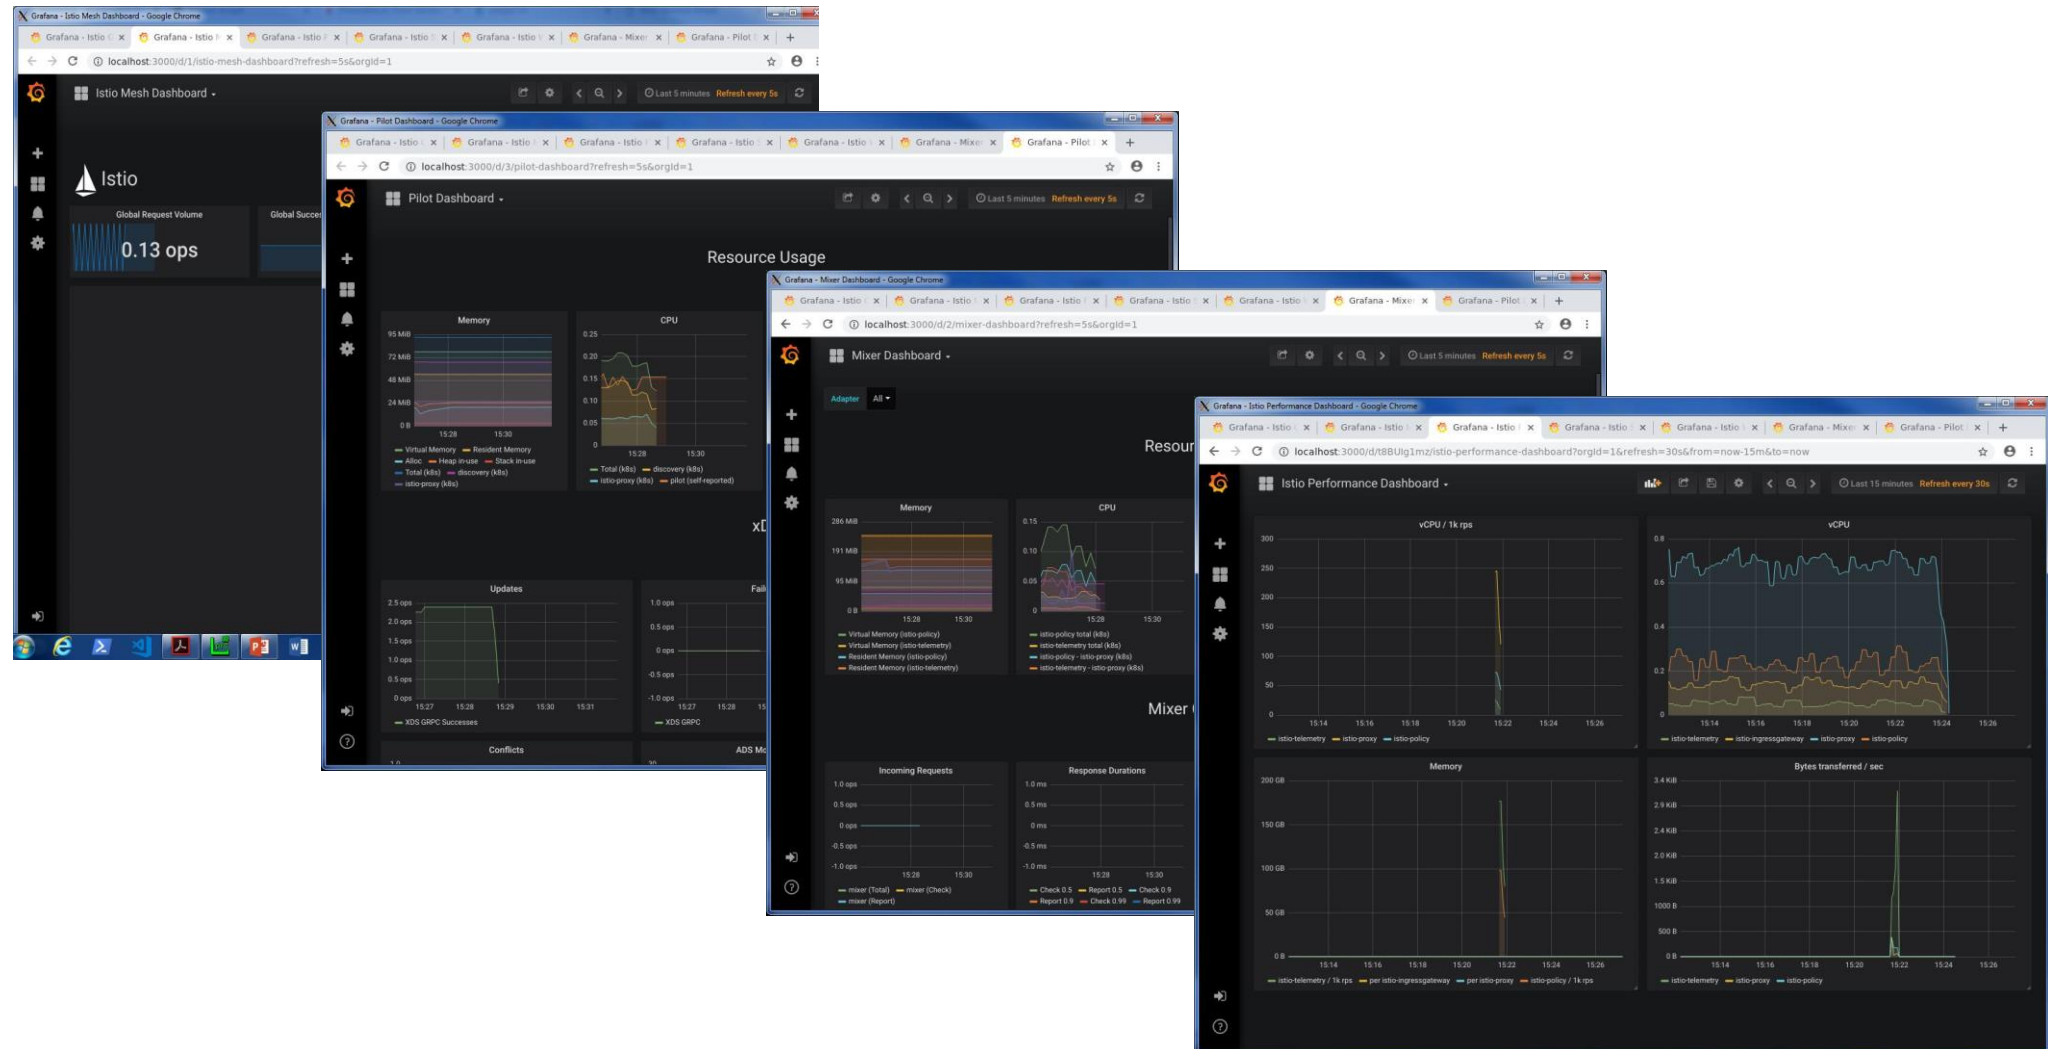
[Querying metrics](#) This task installs the Prometheus add-on for metrics collection and demonstrates querying a configured Prometheus server for Istio metrics.

[Distributed tracing](#) We will now use Istio to trace how requests are flowing across services in the application. Distributed tracing speeds up troubleshooting by allowing developers to quickly understand how different services contribute to the overall end-user perceived latency. In addition, it can be a valuable tool to diagnosis and troubleshooting in distributed applications.

[Using the Istio Dashboard](#) This task installs the Grafana add-on with a preconfigured dashboard for monitoring mesh traffic.

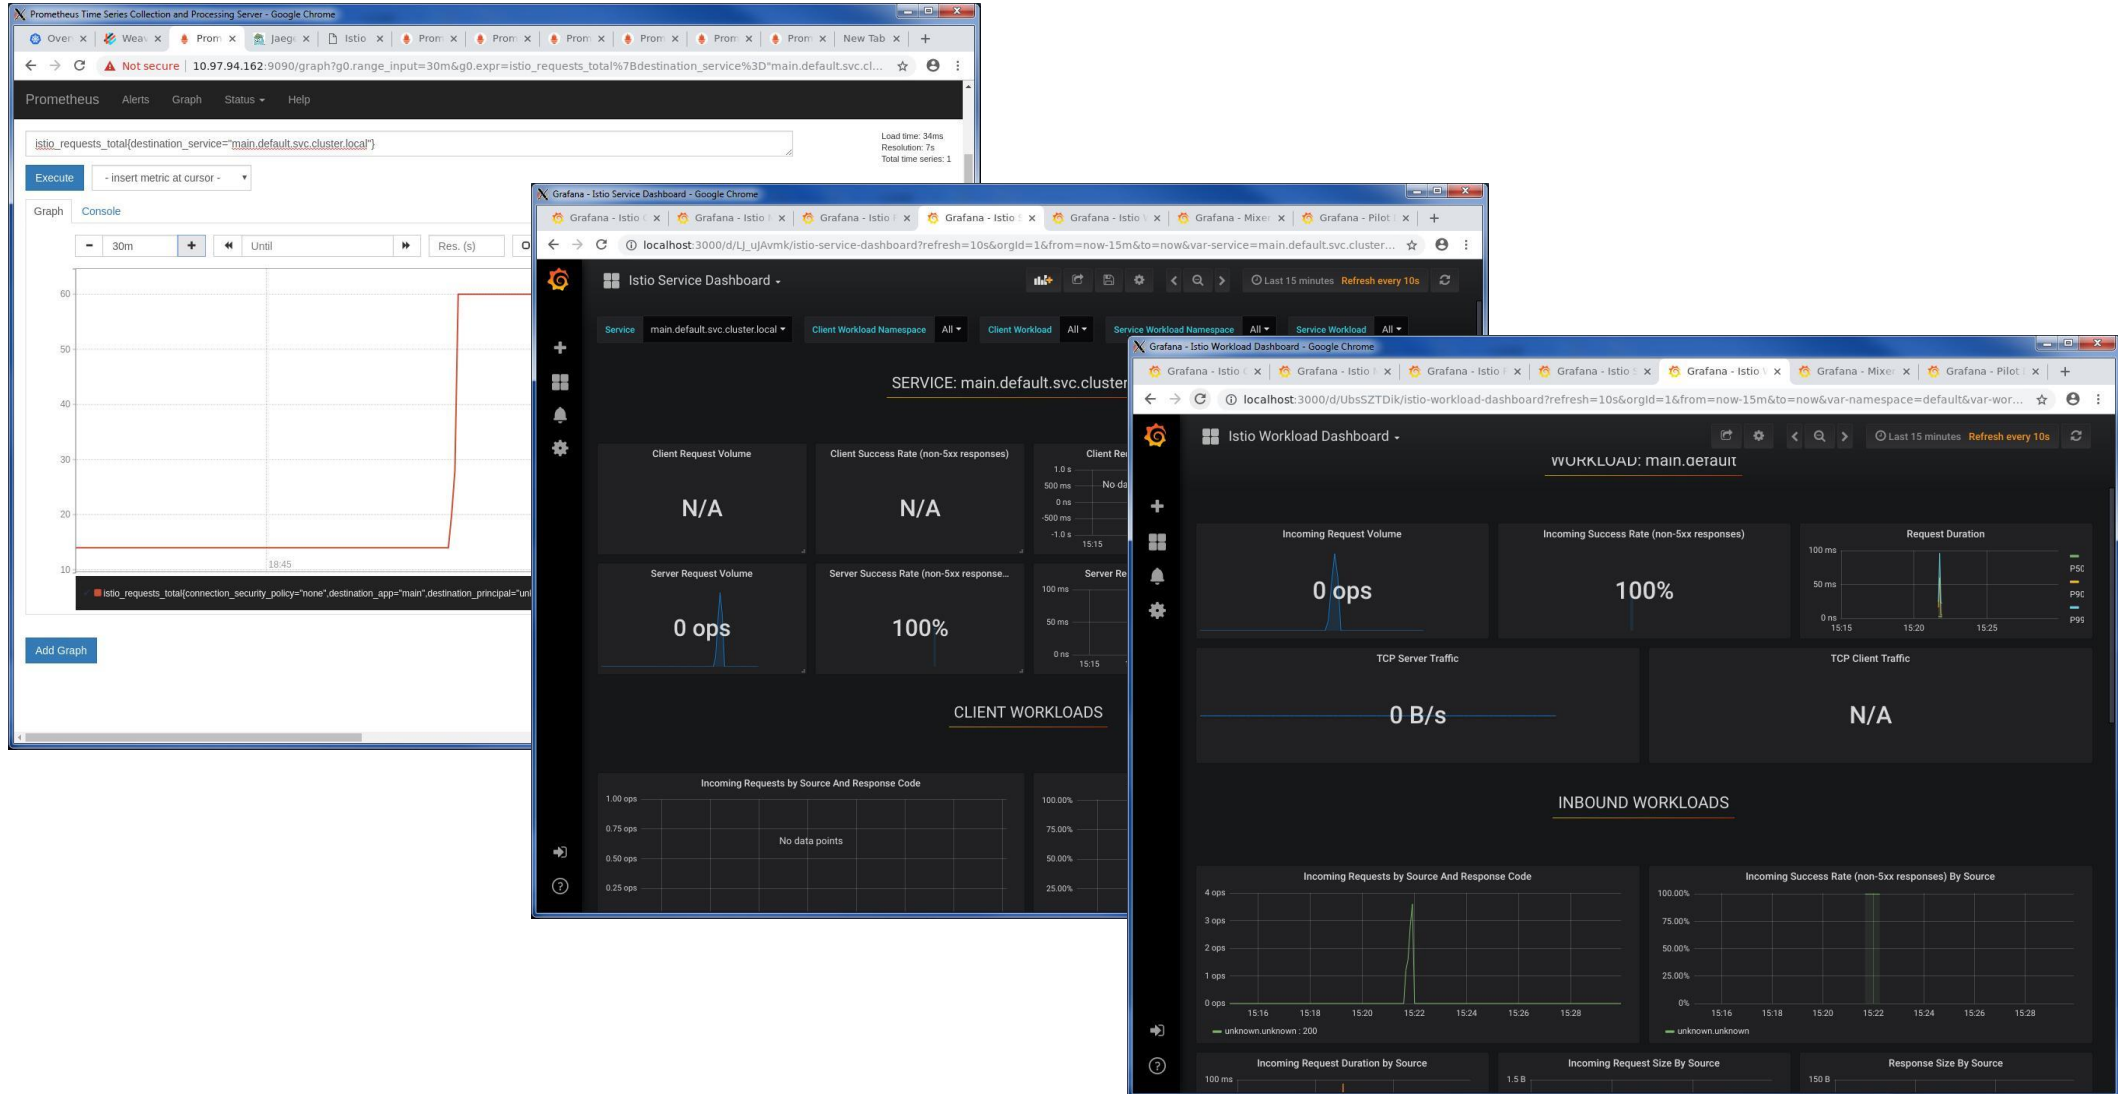
# Prometheus and Grafana (logging/telemetry)

A number of Grafana dashboards are installed with Istio, which track various metrics found in Istio components (like Envoy)



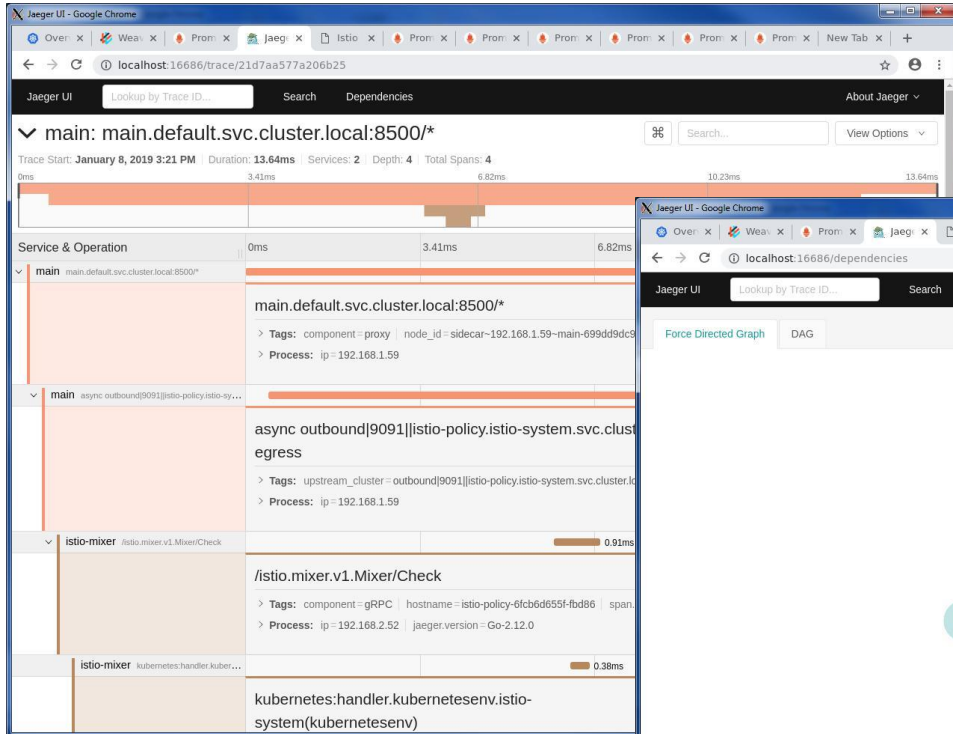
# Prometheus and Grafana (logging/telemetry)

**Main** and **Bookinfo** can be tracked in Prometheus and the Grafana dashboards

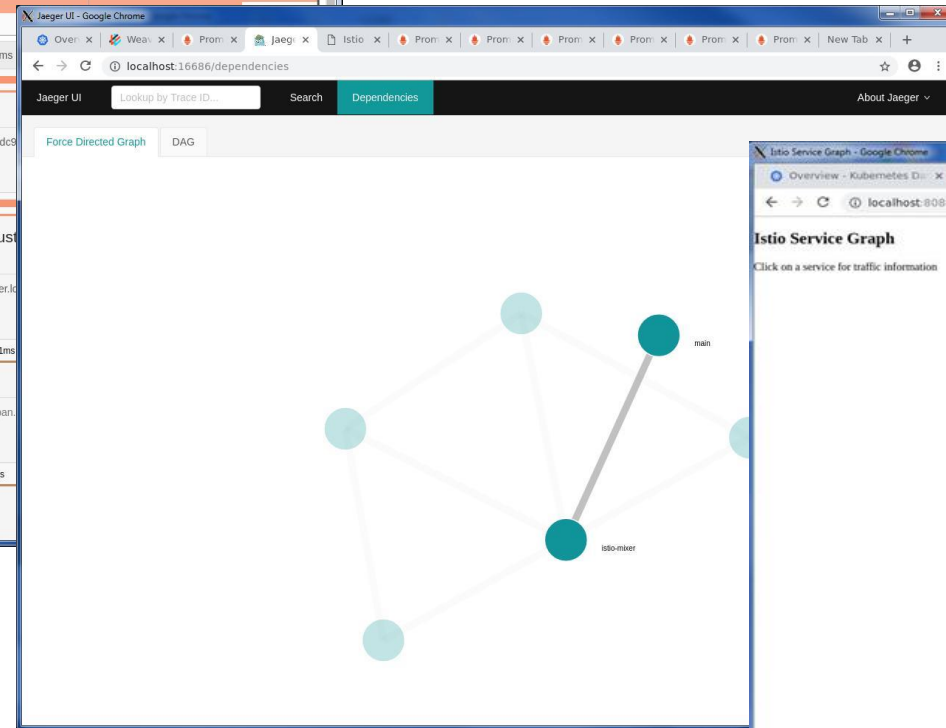


# Distributed/open-tracing (Jaeger) and Service Graph

## Accessing Service *Main*

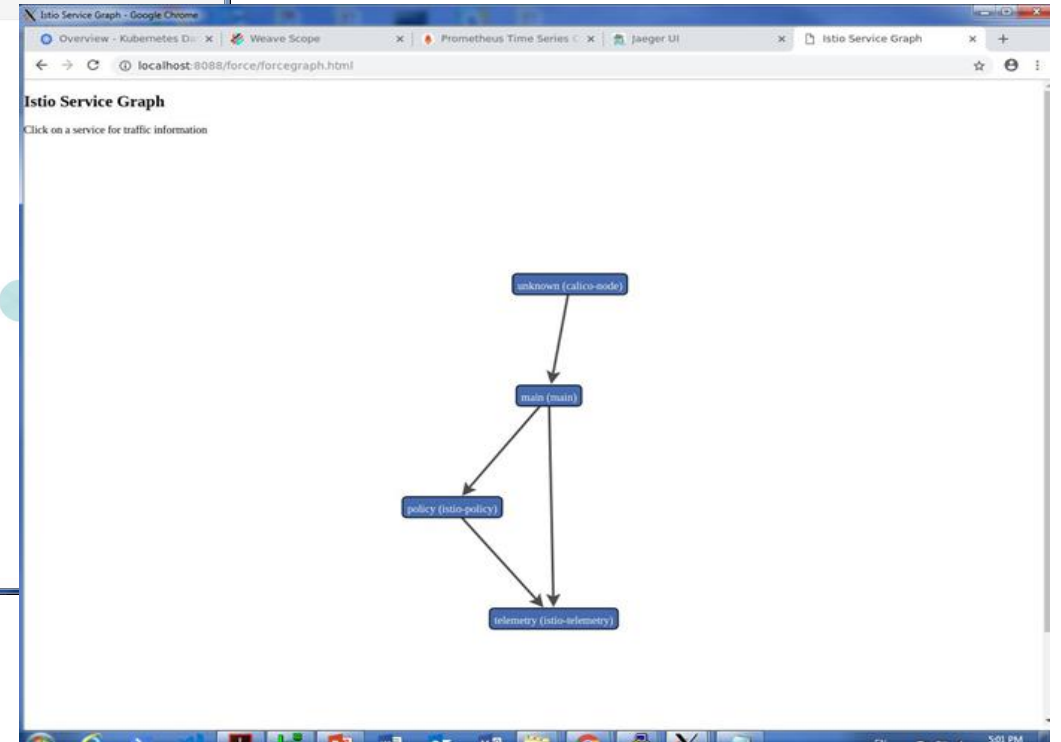


## Tracking Spans, events.... in Jaeger



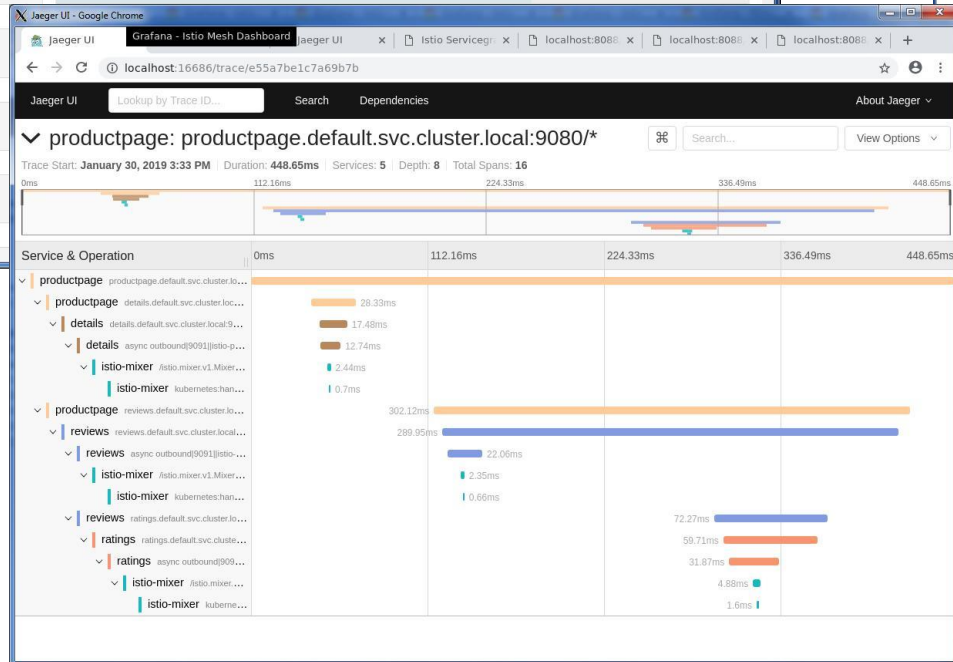
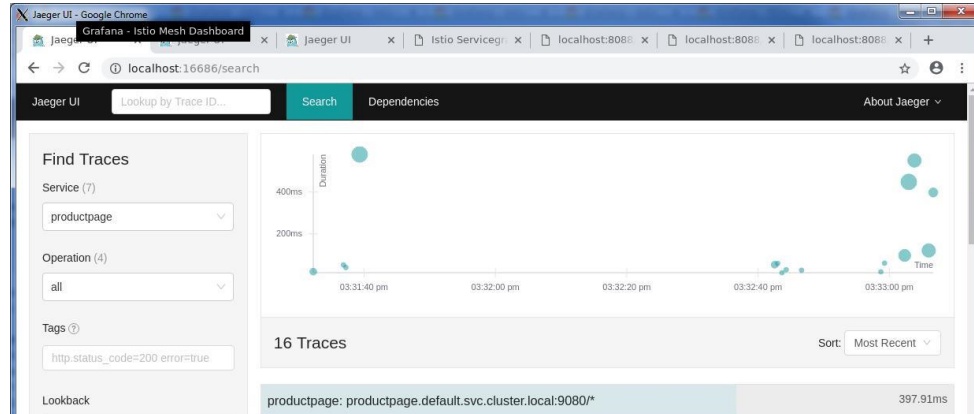
## Force Directed Graph

## Istio Service Graph

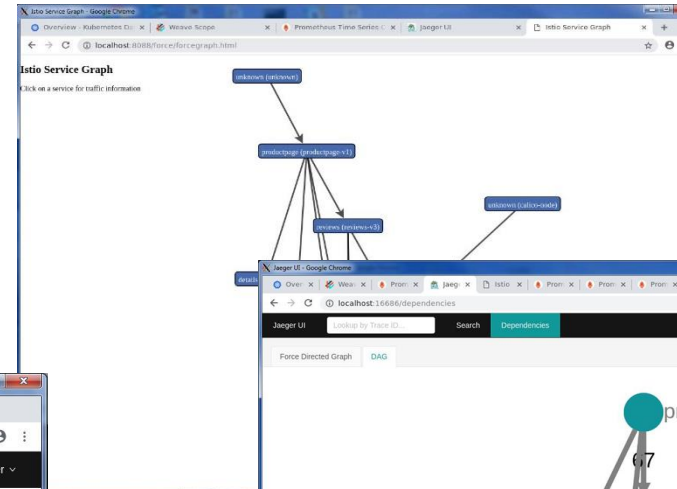


# Distributed/open-tracing (Jaeger) and Service Graph

**Bookinfo** Application – Tracing/graphing comes out of the box as components of Istio

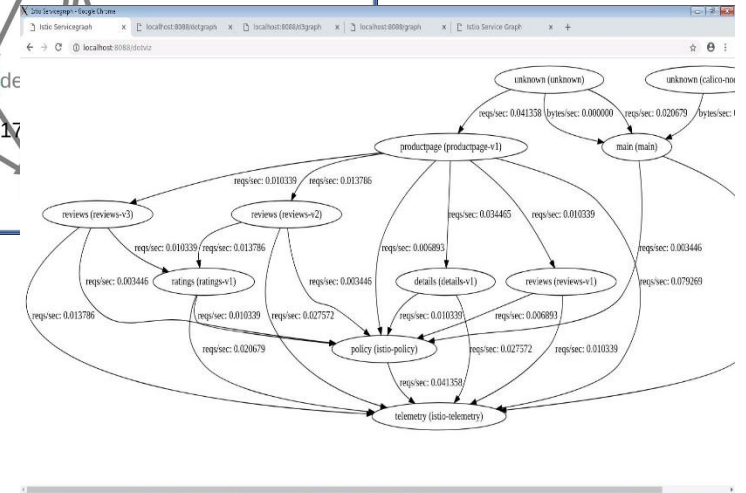


Tracking Spans, events.... in Jaeger



Istio Service Graph

Jaeger directed acyclic graph



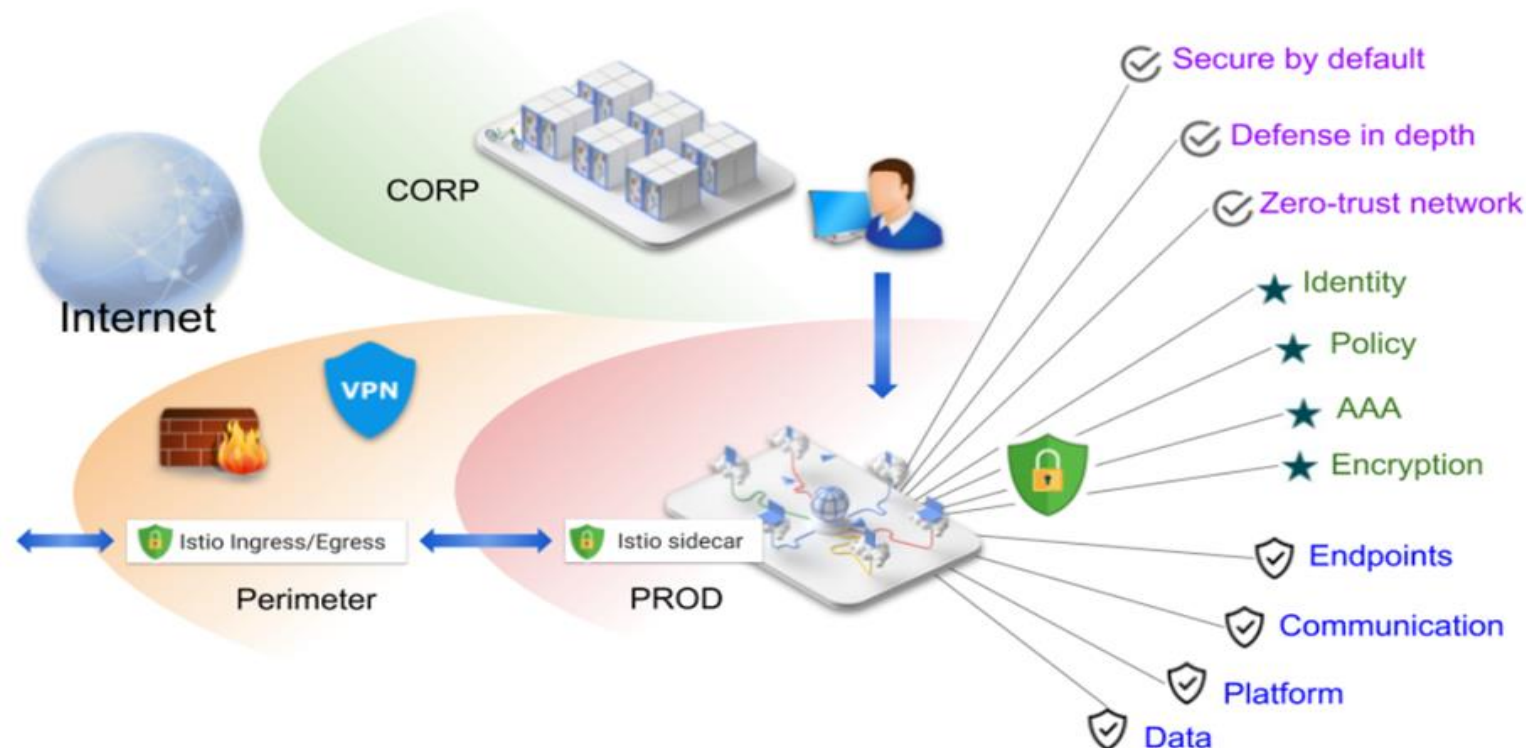
dotviz



# Security

The Istio security features provide strong identity, powerful policy, transparent TLS encryption, and authentication, authorization and audit (AAA) tools to protect your services and data. The goals of Istio security are:

- Security by default: no changes needed for application code and infrastructure
- Defense in depth: integrate with existing security systems to provide multiple layers of defense
- Zero-trust network: build security solutions on untrusted networks



The diagram illustrates the Istio security architecture, showing the flow of data and control/metrics between various components.

**Data Flow (Blue Arrows):**

- Ingress Proxy:** Receives traffic from the Internet (OIDC + TLS / mTLS) and applies **Perimeter security policies**. It connects to **Service A** via **mTLS**.
- Service A:** Contains a **proxy** that routes traffic to **Service B** via **HTTP, gRPC, TCP**.
- Service B:** Contains a **proxy** that applies **Local Authz** and connects to the **Egress Proxy** via **mTLS**.
- Egress Proxy:** Applies **Perimeter security policies** and sends traffic to the Internet (JWT + TLS / mTLS).

**Control + Metrics Flow (Red Dashed Arrows):**

- Pilot:** Manages **Routing + Policy** for the Service A proxy.
- Citadel:** Manages **Cert issuance** for the Service A proxy.
- Mixer:** Manages **Reporting** for the Service B proxy.

**Istio Control Plane:** Consists of **Pilot**, **Citadel**, and **Mixer**, which interact with the **K8S API Server** and the proxies.

**Security Policies:** Can be implemented at different levels of granularity - Service, Namespace, Mesh.

- Security policies can be implemented at different levels of granularity - Service, Namespace, Mesh.

```
graph TD
    Citadel[Citadel] -- CSR --> NodeAgent[Node Agent]
    NodeAgent -- SDS API --> Envoy[Envoy]
    Envoy <--> ServiceFrontend[Service frontend]
    Envoy --- SAN["SAN: 'spiffe://myorg.com/frontend-team'"]
```

The diagram shows the following components and flow:

- Citadel** (orange box with a red pin icon) sends a **CSR** (Certificate Signing Request) to the **Node Agent**.
- Node Agent** (yellow box) communicates with **Envoy** via the **SDS API** (Service Discovery Security API).
- Envoy** (yellow box with a red pin icon) is located within a **Container** on a **K8s Node**.
- Envoy** serves the **Service frontend** (green box).
- The certificate issued to Envoy has a **SAN** (Subject Alternative Name) of `"spiffe://myorg.com/frontend-team"`.



# Conclusion